

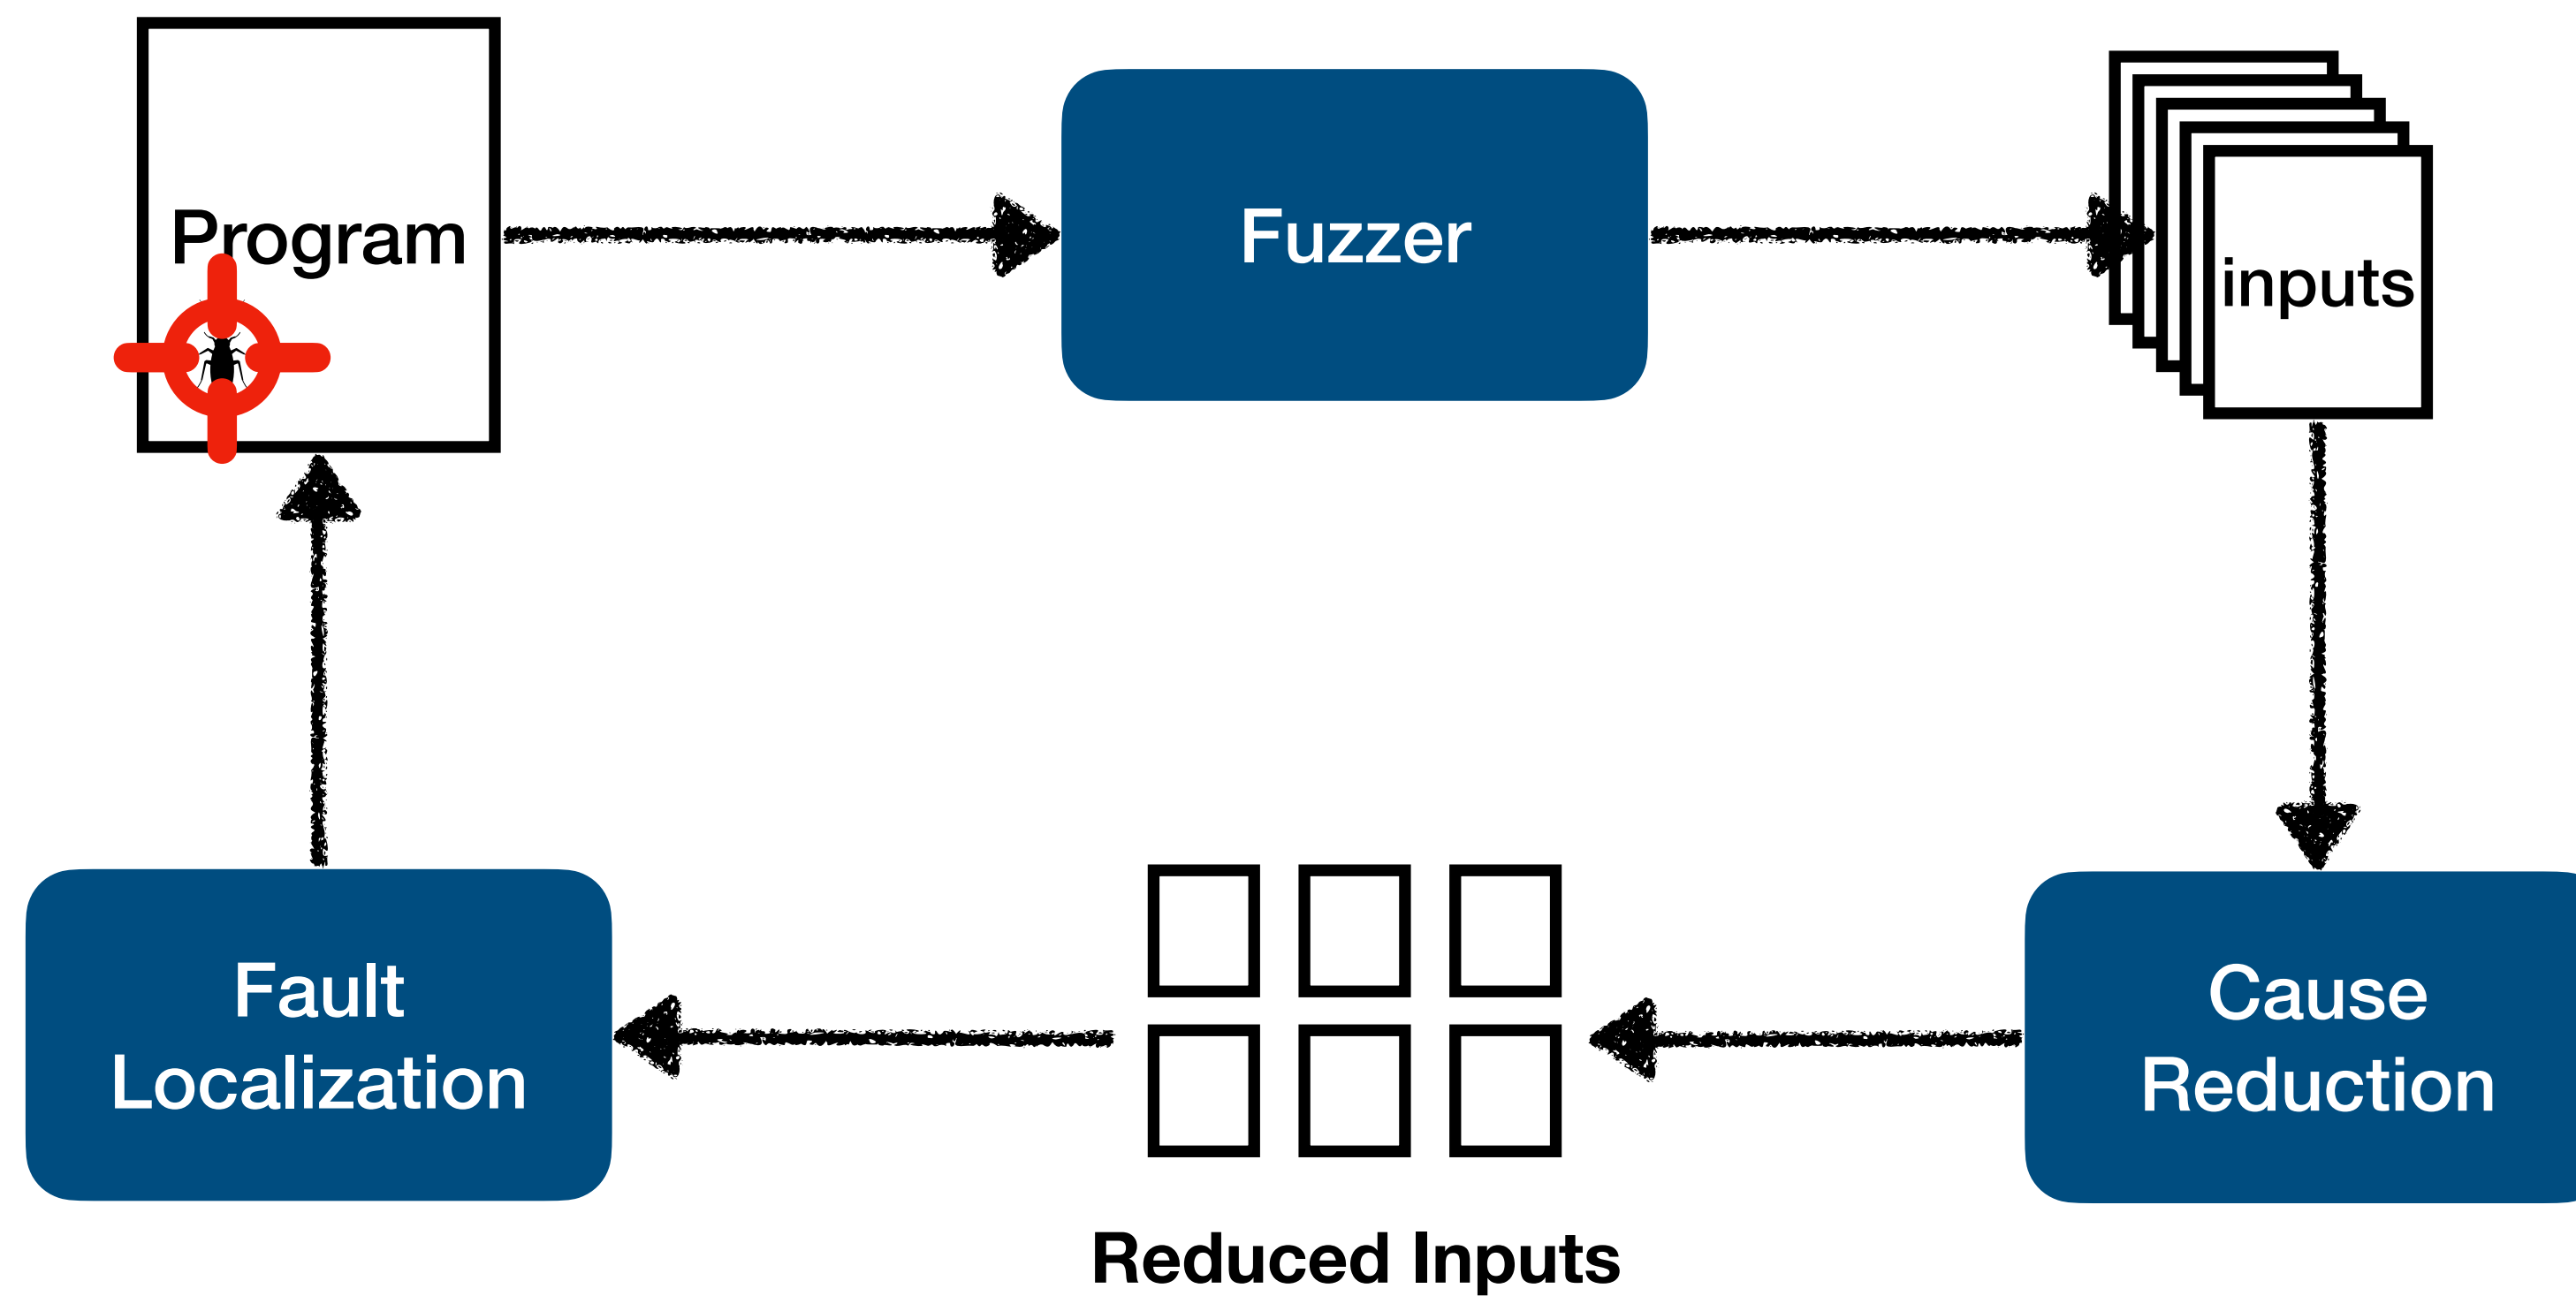
IS893: Advanced Software Security

5. Fault Localization

Kihong Heo



Overview




Fault Localization

- Given 1+ failing inputs and 0+ passing inputs, find a list of suspicious locations
- Two dominant approaches:
 - Statistical: “executing faulty statements is likely to lead to a failure”
 - Logical: “minimal sub-formula making a safety condition UNSAT”

```
1: x = input(); // read “-1”
2: if (x < 0) {
3:   y = x + 1;
4:   z = x + 2;
5: } else {
6:   y = x + 2;
7:   z = x + 2;
8: }
9: k = y;
10: r = z / k; // crash
```

$x = -1 \wedge x < 0 \wedge y_1 = x + 1 \wedge z_1 = x + 2 \wedge (x < 0 ? k = y_1 : k = y_2) \wedge k \neq 0 : \text{UNSAT}$



$x = -1 \wedge y_1 = x + 1 \wedge (x < 0 ? k = y_1 : k = y_2) \wedge k \neq 0 : \text{Minimal UNSAT Core}$

input

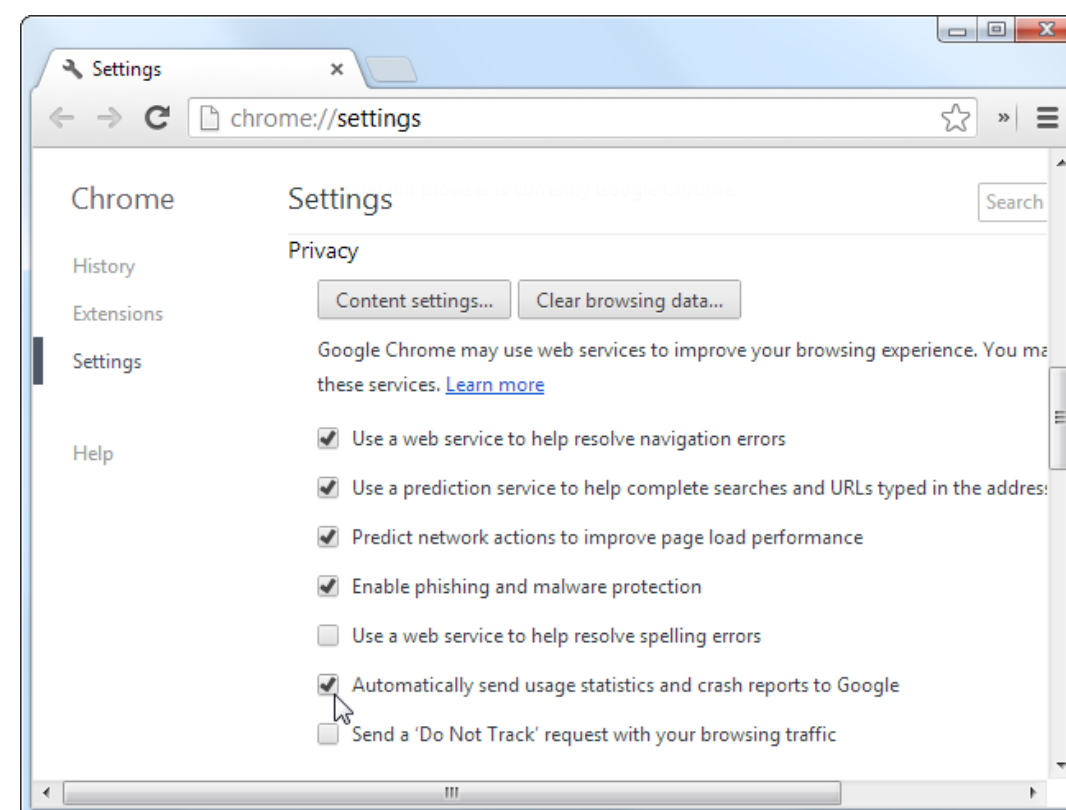
line 3

line 2 & 9

assertion

Statistical Fault Localization

- Statistically analyze correlations between certain behaviors and test results
 - Behavior: code coverage, branch conditions, return values, etc
 - So called “statistical debugging”
 - Enabled by a large amount of test runs from fuzzers, crowdsourcing, etc



Steps

1. **Instrument** the target program to capture “interesting” behaviors
2. **Observe** the behaviors and results for each test run
3. **Analyze** behavioral changes in successful vs. failing runs

Interesting Behavior

- Assume that any interesting behavior is represented as a predicate **P**
- On a program state at a particular program point

Branch Conditions

```
observe(p != 0);  
if (p) {  
    ...  
} else {  
    ...  
}
```

Return Values

```
fd = fopen(...);  
observe(fd > 0);
```

Scala Relationships

```
observe(i < j);  
...  
observe(i > 42);
```

Pointer Relationships

```
observe(p == q);  
...  
observe(p != null);
```

... and many others depending on the problem

Statistical Reasoning

- A large amount of information about many predicates in a program
 - E.g., 300K predicates for UNIX bc with 30K test runs*
- Which predicate is more relevant to bugs?
 - Most of them are not predictive of anything

*Liblit et al., Scalable Statistical Bug Isolation, *PLDI*, 2005

Measures

- How likely is failure when predicate **P** is observed to be true?

$F(P)$ = # failing runs where **P** is observed to be true
 $S(P)$ = # successful runs where **P** is observed to be true

$$\text{Failure}(P) = \text{Pr}(\text{Crash} \mid P \text{ observed to be true}) = \frac{F(P)}{F(P) + S(P)}$$

Hypothesis: **P** is suspicious if $\text{Failure}(P)$ is high

```
f = ...;  
if (f == NULL) {  
    *f;  
}
```


Problem

- “Correlation is not causation”

$F(P)$ = # failing runs where P is observed to be true
 $S(P)$ = # successful runs where P is observed to be true

```
f = ...;  
if (f == NULL) {  
    x = 0;  
    *f;  
}
```

Failure($f == \text{NULL}$) = 1.0
Failure($x == 0$) = 1.0

Fact: High Failure(P) does not mean P is the cause of a bug

More Context

- Intuition:
If **P** is correlated with the failure, regardless of the truth value, then less important
- How likely is failure when predicate **P** is observed?
(not necessarily its truth value)

$F(\mathbf{P} \text{ observed}) = \# \text{ failing runs where } \mathbf{P} \text{ is observed to be true}$ $S(\mathbf{P} \text{ observed}) = \# \text{ successful runs where } \mathbf{P} \text{ is observed to be true}$

$$\text{Context}(\mathbf{P}) = \Pr(\text{Crash} \mid \mathbf{P} \text{ observed to be true}) = \frac{F(\mathbf{P} \text{ observed})}{F(\mathbf{P} \text{ observed}) + S(\mathbf{P} \text{ observed})}$$

New Measure

- How much does **P** being true increase the probability of failure over simply reaching the line where **P** is sampled?

$$\text{Increase(P)} = \text{Failure(P)} - \text{Context(P)}$$

// Suppose 1 failing run and 2 passing runs

```
f = ...;  
if (f == NULL) {  
    x = 0;  
    *f;  
}
```

Failure(f == NULL) = 1.0
Context(f == NULL) = 0.33
Increase(f == NULL) = 0.67

Failure(x == 0) = 1.0
Context(x == 0) = 1.0
Increase(x == 0) = 0.0

“Increase(P) = 1” means high correlation with failing runs

“Increase(P) = 0” means P is an invariant (true for all runs)

“Increase(P) = -1” means high correlation with successful runs

Example

```
void main() {  
    int z;  
    for (int i = 0; i < 3; i++) {  
        char c = getc();  
        if (c == 'a')  
            z = 0;  
        else  
            z = 1;  
        assert(z == 1);  
    }  
}
```

Inputs: {"bba", "bbb"}

Increase	
c == 'a'	0.5
c != 'a'	0.0
i < 3	0.0
i >= 3	-0.5

Simple Algorithm

- Discard predicates having $\text{Increase}(P) \leq 0$
 - E.g., bystander predicates, predicates correlated with success
- Sort remaining predicates by $\text{Increase}(P)$

Real World Example: bc

```
void more_array() {  
    ...  
    /* Copy the old arrays. */  
    for (indx = 1; index < old_count; indx++)  
        arrays[indx] = old_ary[indx];  
    /* Initialize the new elements. */  
    for (; index < v_count; indx++)  
        arrays[indx] = NULL;  
    ...  
}
```

Interesting behavior: scala relationship

Increase	
1	indx > scale
2	indx > use_math
3	indx > opterr
4	indx > next_func
5	indx > i_base
...	...

v_count is a wrong bound

Problem: Multiple Bugs

- Real programs often have multiple unknown bugs
- The effects of bugs may be interrelated
- High Increase values are still good indicators for debugging?
 - Maybe not, because of redundancies
 - E.g., predicates of common bugs may dominates

New Goal


- Find the best predictor for each bug w/o prior knowledge of #bugs
 - Enable us to focus on the bug
- Sorted by the importance of the bugs
 - Enable us to prioritize efforts to the bugs that affect the most users
- How? Intuition: simulate the way humans fix bugs
 - Find the most important bug; fix it; and repeat

New Algorithm

1. Compute various measures (e.g., Increase, F, etc) and rank the predicates
2. Pick the top-ranked predicate **P**
3. Discard all runs where **P** is true:
 - simulate fixing the bug corresponding to **P**
 - reduces rank of correlated predicates:
(predicators of other bugs will rise)
4. Goto 1

Measures

- Is Increase(P) still a good measure?
- Let us organize all measures in compact visual: bug thermometer



Long bar: the predicate was observed many times

Importance

- Ranking by Increase(P)

$$\begin{aligned} \text{Increase(P)} &= \text{Failure(P)} - \text{Context(P)} \\ \text{Failure(P)} &= \text{Pr}(\text{Crash} \mid \text{P observed to be true}) \\ \text{Context(P)} &= \text{Pr}(\text{Crash} \mid \text{P observed}) \end{aligned}$$

(Report from >5000 failing runs)

Thermometer	Context	Increase	S	F	F + S	Predicate
<div></div>	0.065	0.935 ± 0.019	0	23	23	((*(fi + i)))->this.last_token < filesbase
<div></div>	0.065	0.935 ± 0.020	0	10	10	((*(fi + i)))->other.last_line == last
<div></div>	0.071	0.929 ± 0.020	0	18	18	((*(fi + i)))->other.last_line == filesbase
<div></div>	0.073	0.927 ± 0.020	0	10	10	((*(fi + i)))->other.last_line == yy_n_chars
<div></div>	0.071	0.929 ± 0.028	0	19	19	bytes <= filesbase
<div></div>	0.075	0.925 ± 0.022	0	14	14	((*(fi + i)))->other.first_line == 2
<div></div>	0.076	0.924 ± 0.022	0	12	12	((*(fi + i)))->this.first_line < nid
<div></div>	0.077	0.923 ± 0.023	0	10	10	((*(fi + i)))->other.last_line == yy_init
..... 2732 additional predictors follow						

High Increase(P) scores tend to report predicates with high Failure and low Context
=> if observed to be true, very likely crash, but mostly observed to be false
=> indicate very special cases of more general bugs

Importance

- Ranking by F(P)


$$\begin{aligned} \text{Increase(P)} &= \text{Failure(P)} - \text{Context(P)} \\ \text{Failure(P)} &= \text{Pr}(\text{Crash} \mid \text{P observed to be true}) \\ \text{Context(P)} &= \text{Pr}(\text{Crash} \mid \text{P observed}) \end{aligned}$$

(Report from >5000 failing runs)

Thermometer	Context	Increase	S	F	F + S	Predicate
<div><div></div></div>	0.176	0.007 ± 0.012	22554	5045	27599	files[filesindex].language != 15
<div><div></div></div>	0.176	0.007 ± 0.012	22566	5045	27611	tmp == 0 is FALSE
<div><div></div></div>	0.176	0.007 ± 0.012	22571	5045	27616	strcmp != 0
<div><div></div></div>	0.176	0.007 ± 0.013	18894	4251	23145	tmp == 0 is FALSE
<div><div></div></div>	0.176	0.007 ± 0.013	18885	4240	23125	files[filesindex].language != 14
<div><div></div></div>	0.176	0.008 ± 0.013	17757	4007	21764	filesindex >= 25
<div><div></div></div>	0.177	0.008 ± 0.014	16453	3731	20184	new value of M < old value of M
<div><div></div></div>	0.176	0.261 ± 0.023	4800	3716	8516	config.winnowing_window_size != argc
..... 2732 additional predictors follow						

High F(P) scores tend to report predicates with low Failure and high Context
=> if observed to be true, more likely pass
=> indicate very general predicates covering many different bugs as well as correct behaviors









Analogy



Harmonic Mean

- Standard solution to achieve both high precision and high recall

$$\text{Importance(P)} = \frac{2}{1/\text{Increase(P)} + 1/\text{F(P)}}$$

Thermometer	Context	Increase	S	F	F + S	Predicate
	0.176	0.824 ± 0.009	0	1585	1585	files[filesindex].language > 16
	0.176	0.824 ± 0.009	0	1584	1584	strcmp > 0
	0.176	0.824 ± 0.009	0	1580	1580	strcmp == 0
	0.176	0.824 ± 0.009	0	1577	1577	files[filesindex].language == 17
	0.176	0.824 ± 0.009	0	1576	1576	tmp == 0 is TRUE
	0.176	0.824 ± 0.009	0	1573	1573	strcmp > 0
	0.116	0.883 ± 0.012	1	774	775	((*(fi + i)))->this.last_line == 1
	0.116	0.883 ± 0.012	1	776	777	((*(fi + i)))->other.last_line == yyleng
..... 2732 additional predictors follow						

Conclusion

- Which part of the program is wrong? Fault localization!
- A common approach: statistically localize root causes of bugs
 - Collect data by instrumentation and crash reports (fuzzing, crowd, etc)
 - Many metrics to rank “interesting behaviors”
- Applications: guiding manual debugging or automated program repair tools